

# Module 1

---

## Exercise 1: A look at PostgreSQL pages

---

1. Connect to your instance with **pgadmin** and create a database
2. Create a table containing only one **VARCHAR** column, insert 'Hello World!' in this table.
  - 💡 do not forget to **CHECKPOINT**!
  - What did this achieve?
3. Query the **pg\_database** table to find the OID (Object Identifier) of your database
4. Exec into your database container. Go to **/var/lib/postgresql/data/base/** and find the folder containing your database. What is inside?
5. Find the file for your table using **pg\_relation\_filepath()**. Use the **du -sb** command to check its size in bits. What is this size? Why?
6. Install **hexdump** in the container using **apt update && apt install bsdmainutils**
7. Run **hexdump -C** on your table file. What do you see?
8. Insert another value in the table. Check again.
  - 💡 do not forget to **CHECKPOINT**!
9. Same question with a **NULL** value
10. Now delete the whole table and check again. What is happening? How could you fix this?

## Exercise 2: The Write-Ahead Log (WAL)

---

1. Connect to your instance with **pgadmin** and create a database or use an existing one
2. Install the extension **pg\_walinspect** with **CREATE EXTENSION**
3. Create a table of any structure
4. Write down the current LSN with **SELECT pg\_current\_wal\_lsn()**
5. Interact with your table: insert a few lines for example.
6. Note the new LSN with **SELECT pg\_current\_wal\_lsn()**
7. Use the **pg\_get\_wal\_records\_info()** function to get all the WAL records between the two LSNs you noted.
8. Try different operations (insertions, deletions, updates) and look at the WAL content.

## Questions

- Describe the output of the **pg\_get\_wal\_records\_info()** function. What information does it extract from the WAL?

## Exercise 3: Transactions Isolation

---

1. Create a simple table named **test** with two INT columns named **id** and **value**
2. Open two tabs in **pgadmin**. We will call them **tab1** and **tab2** from now on
3. In **tab1**, run **BEGIN** to start a transaction. In this same tab, run an **INSERT** command to add a new line. Finally run a **SELECT \*** from the table in both **tab1** and **tab2**.

- What do you see?
  - Why is that? What is the isolation level you are using?
4. **COMMIT** the transaction in tab1. Run again the **SELECT** in both tabs. What happened?
  5. Now in tab1 start a new transaction with **BEGIN TRANSACTION ISOLATION LEVEL REPEATABLE READ**;
  6. **SELECT** the content of the table. Go to the second tab and add a new record. Go back to the first one and **SELECT** again.
- What do you see?
7. **ROLLBACK** the transaction in tab1.
  8. Finally now use **BEGIN TRANSACTION ISOLATION LEVEL SERIALIZABLE**; in tab1. Run an update on one of the lines in tab2. Now run the same exact UPDATE in tab1. What is happening?

## Exercise 4: Explicit locks

---

1. Create a new table of any form. Insert a few records in this table.
2. Start a transaction in a tab with **BEGIN**. Issue a

```
SELECT *  
FROM <table>  
FOR UPDATE
```

3. Go to a second tab. Run the following:

```
SELECT l.*  
FROM pg_locks l  
JOIN pg_class t ON l.relation = t.oid  
WHERE t.relkind = 'r'  
      AND t.relname = 'toto';
```

- What do you see?
4. Now in the same tab, try to modify the table, for example run an update? What happens?
  5. Go back to the first tab and **ROLLBACK** the transaction. What happened in the second tab? Go to <https://www.postgresql.org/docs/current/explicit-locking.html> to understand a bit more about what you just did.