


Module 1

Storing Data Safely

Numbers any engineer should know



Pros and cons of disks

👍 They have a bigger capacity than volatile memory or caches

👍 They are **persistent**

😞 They are **slow**

How do this?

 Any guess?

Using the OS

`mmap` allows for mapping in virtual memory a physical file

But this is problematic. Why? 🙋

Using the OS

`mmap` allows for mapping in virtual memory a physical file

But this is problematic. Why? 🙋

For **performance** reasons mostly. Pages can be evicted at any time, which could stall processes

Buffer Pool

Most DBMS choose instead to, implement their own layer, sitting **between** the database process and the filesystem.

Buffer Pool

Database files are split in **pages** of a few KBs

A page is the **smallest** unit of storage the DBMS can interact with

➡ Postgres uses 8KB pages by default (same for Microsoft SQL Server)

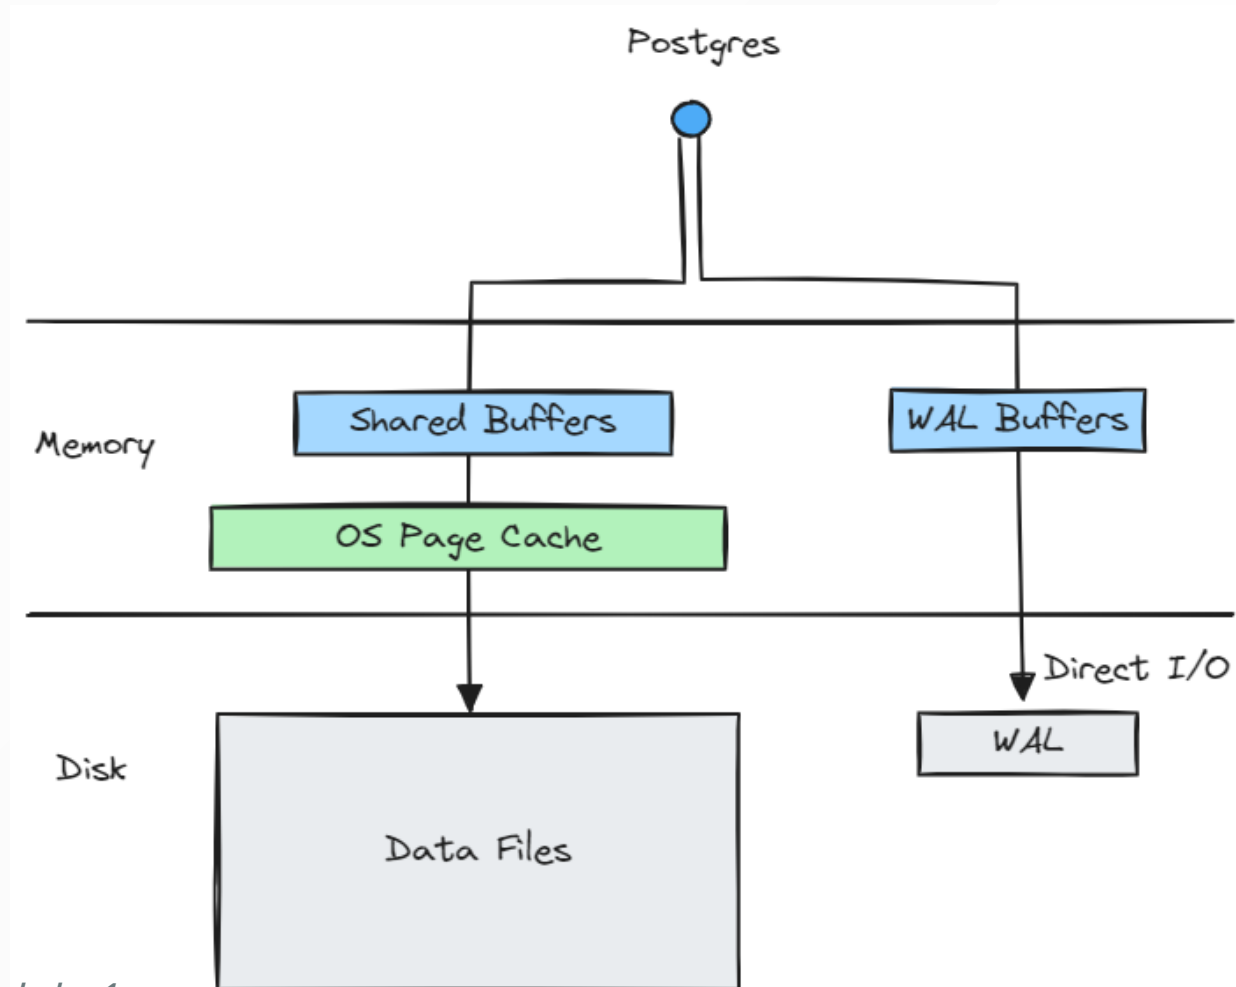
The **Buffer Pool** is a cache of pages, organized as an in-memory array


Buffer Pool

The Buffer decides **when** data is read from disk to memory, and **when** pages are persisted to disk

The Buffer Pool generally does `O_DIRECT`, but Posgres does not.

Anatomy of Postgres storage





Postgres uses Buffered I/O ?

It started as a **research** project in the 1980s and only has still a few contributors

Main focus was on the **database engine** itself

A Direct I/O stack is **hard to implement**

Data Files

How are **standard** types stored?

 How are **NULL**s stored?

For big enough values, store them in specific pages (**TOAST** pages in Postgres)

TOAST = **T**he **O**versized-**A**tttribute **S**torage **T**echnique

Tuple Alignment

Alignment: Making sure that tuples are stored on multiples of 64 bits.

Some systems reorganize columns in the tuple.

PostgreSQL will **pad** every type to make sure that everything is **64-bit** aligned

Holy grail of linearizability

Database systems would be **easier to write** if they were single process and single user

We would love transactions to appear as if they were executed in a **global order**, even if they are executed **concurrently**

This property is called **linearizability** and is often used in distributed systems

Dealing with concurrency

However in real life we will have **multiple users** running multiple concurrent queries

How do we solve this?

Transaction

A sequence of **Read** and **Write** operations that has the four well known **ACID** properties



Remember what they stand for?

ACID

Atomicity: All changes commit or none

Consistency: They only transition from a valid state to another

Isolation: They execute isolated from each other

Durability: Committed transactions persist despite system failures

Atomicity & Durability

To provide **both**, the best way is to use a **log**

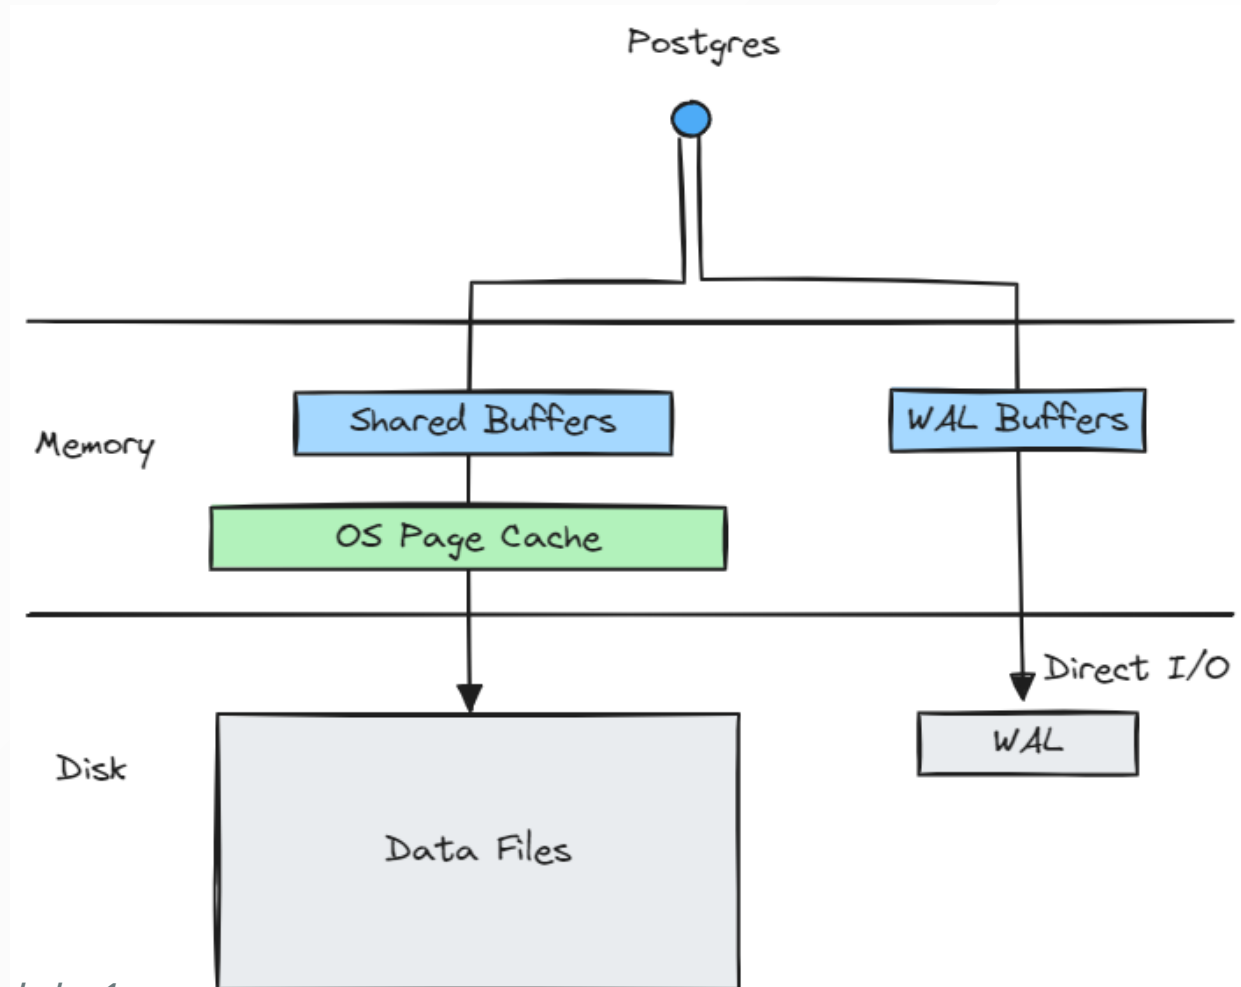
First write all operations in a queue, and apply them once committed

Write-Ahead Log (WAL)

All **operations** are written in the log

CHECKPOINT is a special one

Recap



Checkpoint

Happens automatically **every few minutes**

1. Get the current WAL **position**
 2. Write data from Buffer Pool to **page cache**
 3. Call `fsync` on **all modified files**
 4. Mark the point of recovery, and clean the log
- ? Did you recently hear about **issues** there (think of `fsync`)?

Isolation Anomalies

Isolation is another story

We are trying to prevent the following to happen:

- **Dirty** reads: reading uncommitted changes
- **Non-repeatable** reads: reading a value has been changed by another committed transaction
- **Phantom Reads** reads: reading a value that has been inserted by another committed transaction

Isolation Levels

- **Read Uncommitted:** Allows dirty reads, non-repeatable reads, and phantom reads
- **Read Committed:** Prevents dirty reads but allows non-repeatable and phantom reads
- **Repeatable Read:** Prevents dirty and non-repeatable reads but allows phantom reads
- **Serializable:** Ensures the highest level of isolation, preventing all anomalies

How to guarantee this?

There are mainly two approaches.

 Can you name one?

How to guarantee this?

There are mainly three approaches.

 Can you name one?

MVCC, Locks and Optimistic CC

Optimistic Concurrency Control

Assume that everything *generally* goes fine

1. Record the transaction beginning timestamp
2. Prepare the values to write
3. If the destination has not changed since the beginning of the transaction, **COMMIT**, otherwise, **ROLLBACK**

Locks

A **Lock Manager** grants different kinds of locks on objects, such as **Exclusive** or **Shared** locks

Some locks are compatible with each other.

in most DBMS, `SELECT` will acquire a S lock that is compatible with other S locks. However, modifying the table requires an X lock

Locks limit the **concurrency**

MVCC


Every operation operates on a **snapshot** of the database

Reading never blocks writing and writing never blocks reading.

Tuples are not modified in **place**

Problems: outdated blocks need to be **cleaned up** and txid wraparound 🦴

Isolation Levels in PostgreSQL

- `REPEATABLE_READ` is provided by taking a snapshot at the beginning of the transaction. Concurrent updates will lead to a so-called serialization error and a `ROLLBACK`
- `SERIALIZABLE` is the same with a stronger guarantee: if a table is **read** by two concurrent transactions, one is aborted
- `READ_COMMITTED` is the default
- `READ_UNCOMMITTED` does not really exist
 Why?

Recap

- We looked at how PostgreSQL **stores data on disk**
- We saw that it uses a **log** named the **WAL** to ensure atomicity and durability
- Isolation is provided by **transactions** that have different isolation levels